

RNN\_assignmnet.ipynb - Colaboratory

Name: Niraj Lamichhane

Group Leader: Niraj Lamichhane

Group Members:Rojan Shrestha,Saugat Karki, Niraj Lammichane, Aayush Niraula

Group: L6CG5

University\_id: 2059514

Module Leader: Mr.Siman Giri

Tutor: Mr.Akash Adhikari

```
from google.colab import drive # load drive and mount
drive.mount('/content/drive')

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

import pandas as pd
import re
from sklearn.model_selection import train_test_split
from keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from keras.models import Sequential
from keras.layers import Embedding, LSTM, Dense
from keras.callbacks import EarlyStopping
import matplotlib.pyplot as plt
import tensorflow as tf
from sklearn.metrics import confusion_matrix, f1_score, roc_curve, auc
```

```
# load the dataset
data_path = "/content/drive/MyDrive/Amazon Book Review/kindle_review.csv"
df = pd.read_csv(data_path)

# print the first five rows of the dataset
print(df.head())
```

	Unnamed: 0	rating	reviewText
0	0	5	This book was the very first bookmobile book I...
1	1	1	When I read the description for this book, I c...
2	2	5	I just had to edit this review. This book is a...
3	3	5	I don't normally buy 'mystery' novels because ...
4	4	5	This isn't the kind of book I normally read, a...

	summary
0	50+ years ago...
1	Boring! Boring! Boring!
2	Wiggleslicious/new toy ready!!

```
3          Very good read.
4          Great Story!
```

```
df.isnull().any() # null value checked
```

```
Unnamed: 0    False
rating        False
reviewText    False
summary       False
dtype: bool
```

```
# perform basic cleaning tasks
df = df[['reviewText', 'rating']] # select review and rating only
df = df.dropna() # remove the rows with missing values
df = df[df['rating'] != 3] # remove rows where the rating is 3
df['sentiment'] = df['rating'].apply(lambda x: 1 if x > 3 else 0) ## Create a new 'sentiment' column based on the 'rating' column,
df = df.drop(columns=['rating']) # Drop the 'rating' column, since we no longer need it
```

	reviewText	rating
0	This book was the very first bookmobile book I...	5
1	When I read the description for this book, I c...	1
2	I just had to edit this review. This book is a...	5
3	I don't normally buy 'mystery' novels because ...	5
4	This isn't the kind of book I normally read, a...	5

```
# remove unwanted text and characters
df['reviewText'] = df['reviewText'].apply(lambda x: re.sub('[^a-zA-Z0-9\s]', '', x))
```

```
def remove_urls(text_series):
    # Define a regular expression pattern to match URLs or website addresses
    url_pattern = re.compile(r'https?://\S+|www\.\S+')
    # Use the apply() method to apply the regular expression pattern to each string in the input series, and replace any matches with an empty string
    return text_series.apply(lambda x: url_pattern.sub('', x))
```

```
text = df['reviewText']
# Pass the 'text' variable containing the 'reviewText' column of the DataFrame to the 'remove_urls' function to remove any URLs or website addresses
remove_urls(text)
remove_urls(text)
```

```
0      This book was the very first bookmobile book I...
1      When I read the description for this book I co...
2      I just had to edit this review This book is an...
3      I dont normally buy mystery novels because I j...
4      This isnt the kind of book I normally read alt...
...
...
11994  After E A Poe came H P Lovecraft in the world ...
11995  Had to read certain passages twicetypos Wish ...
11997  Dragon Knights is a world where Knights ride d...
11998  Since this story is very short its hard to say...
11999  from 1922 an amazing collection of info on sym...
Name: reviewText, Length: 10000, dtype: object
```

```
# normalize text data
df['reviewText'] = df['reviewText'].apply(lambda x: x.lower())
```

## ▼ Build Model[5]:

```
# split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(df['reviewText'], df['sentiment'], test_size=0.2, random_state=42)
```

```
# tokenize and pad the sequences
tokenizer = Tokenizer(num_words=5000, oov_token='UNK') # Create a tokenizer object
tokenizer.fit_on_texts(X_train) # Fit the tokenizer on the training data to generate a vocabulary of unique words
```

5/8/23, 6:27 PM

RNN\_assignment.ipynb - Colaboratory

```
X_train_seq = tokenizer.texts_to_sequences(X_train)# Convert each text document in the training data to a sequence of integers using the token
X_test_seq = tokenizer.texts_to_sequences(X_test)# Convert each text document in the test data to a sequence of integers
X_train_padded = pad_sequences(X_train_seq, maxlen=200) # Pad the sequences of integers with zeros to ensure that they all have the same length
X_test_padded = pad_sequences(X_test_seq, maxlen=200) # This is necessary because neural networks expect inputs with a fixed size
```

```
# print the shape of the cleaned and preprocessed dataset
print(X_train_padded.shape)
print(X_test_padded.shape)

(8000, 200)
(2000, 200)
```

```
# define the model architecture
model = Sequential()
model.add(Embedding(input_dim=5000, output_dim=32, input_length=200))
model.add(LSTM(units=64, dropout=0.2, recurrent_dropout=0.2))
model.add(Dense(units=1, activation='sigmoid'))

# compile the model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

# print the model summary
print(model.summary())
```

```
Model: "sequential"
_____
Layer (type)                 Output Shape              Param #
-----
embedding (Embedding)        (None, 200, 32)          160000
lstm (LSTM)                   (None, 64)                24832
dense (Dense)                 (None, 1)                 65
-----
Total params: 184,897
Trainable params: 184,897
Non-trainable params: 0
_____
None
```

## ▼ Training of the Model[5]:

```
# Define a custom callback class


class Mycallback(tf.keras.callbacks.Callback):
    def on_epoch_end(self, epoch, logs={}): # Define the 'on_epoch_end' method, which will be called by Keras after each training epoch
        if(logs.get("accuracy")>0.95): # Check if the accuracy of the model is greater than 0.95
            print('\nloss is low so stop training')# Print a message indicating that the training will stop because the loss is low
            self.model.stop_training = True# Set the 'stop_training' attribute of the model to True to stop training

callbacks=Mycallback() #This code creates an instance of the Mycallback class

# train the model
history = model.fit(X_train_padded, y_train, epochs=10, batch_size=128, validation_data=(X_test_padded, y_test), callbacks=[callbacks])

# plot the training and validation loss
# Plot the training and validation loss over time using Matplotlib
# The 'history' object returned by the 'fit()' method contains information about the training process
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.legend()
plt.show()
```

Epoch 1/10  
63/63 [=====] - 38s 533ms/step - loss: 0.6379 - accuracy: 0.6404 - val\_loss: 0.5650 - val\_accuracy: 0.7715  
Epoch 2/10  
63/63 [=====] - 31s 496ms/step - loss: 0.4452 - accuracy: 0.8240 - val\_loss: 0.3490 - val\_accuracy: 0.8650  
Epoch 3/10  
63/63 [=====] - 31s 498ms/step - loss: 0.2789 - accuracy: 0.8880 - val\_loss: 0.2982 - val\_accuracy: 0.8800  
Epoch 4/10  
63/63 [=====] - 31s 498ms/step - loss: 0.2021 - accuracy: 0.9247 - val\_loss: 0.2598 - val\_accuracy: 0.8965  
Epoch 5/10  
63/63 [=====] - 33s 532ms/step - loss: 0.1586 - accuracy: 0.9424 - val\_loss: 0.2767 - val\_accuracy: 0.9010  
Epoch 6/10  
63/63 [=====] - ETA: 0s - loss: 0.1277 - accuracy: 0.9570  
Loss is low so stop training  
63/63 [=====] - 31s 500ms/step - loss: 0.1277 - accuracy: 0.9570 - val\_loss: 0.2720 - val\_accuracy: 0.9000



## ▼ Evaluate the model[2.5]:

```
# evaluate the model
# Evaluate the trained model on the test data using the evaluate() method
# The 'X_test_padded' and 'y_test' data are used for evaluation
# The 'verbose' argument controls the verbosity of the evaluation output
loss, accuracy = model.evaluate(X_test_padded, y_test, verbose=0)
# Print the test loss and accuracy
print('Test Loss:', loss)
print('Test Accuracy:', accuracy)
```

```
Test Loss: 0.27198123931884766
Test Accuracy: 0.8999999761581421
```

```
# get the predicted labels
y_pred = model.predict(X_test_padded)
y_pred = [round(pred[0]) for pred in y_pred]

# create the confusion matrix
cm = confusion_matrix(y_test, y_pred)

# print the confusion matrix
print(cm)
```

```
63/63 [=====] - 5s 80ms/step
[[ 705  124]
 [  76 1095]]
```

```
# get the predicted probabilities
y_prob = model.predict(X_test_padded)

# get the predicted labels
y_pred = [round(prob[0]) for prob in y_prob]

# calculate the F1 score
f1 = f1_score(y_test, y_pred)


# print the F1 score
print('F1 score:', f1)
```

```
63/63 [=====] - 2s 34ms/step
F1 score: 0.9163179916317991
```

```
# calculate the ROC curve
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
```

```
# calculate the AUC score
auc_score = auc(fpr, tpr)

# plot the ROC curve
plt.plot(fpr, tpr, label='ROC curve (AUC = %0.2f)' % auc_score)
plt.plot([0, 1], [0, 1], linestyle='--', label='Random guess')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```



## ▼ Results and Prediction[2.5]:

```
# make predictions on the test set
y_pred = model.predict(X_test_padded)
y_pred = [round(pred[0]) for pred in y_pred]

# print the first 10 true and predicted labels
print('True labels:', list(y_test)[:10])
print('Predicted labels:', y_pred[:10])
```

```
63/63 [=====] - 2s 35ms/step
True labels: [1, 1, 1, 0, 1, 0, 0, 0, 1, 1]
Predicted labels: [1, 1, 1, 0, 1, 0, 1, 0, 1, 0]
```

```
neg_review = ["I bought this set and returned it. I couldn't force myself to finish the first book, A Touch of Silk. ( They are bundled out
```

```
# Let's tokenize it and do the pad_sequence to make it in right format acceptable by model
neg_review_token = tokenizer.texts_to_sequences(neg_review)
```

```
# padding
neg_review_padded = pad_sequences(neg_review_token, maxlen=100, padding='post')
```

```
review_predict = (model.predict(neg_review_padded)>0.5).astype('int32')
```

```
1/1 [=====] - 1s 613ms/step
```

```
# 1 is Positive review and 0 is negative review
if review_predict[0] == 0:
    print("It's a negative review")
else:
    print("It's a positive review")
```

```
It's a negative review
```

```
pos_review = ["This book was the very first bookmobile book I bought when I was in the school book club. I loved the story then and I bet a d
```

5/8/23, 6:27 PM

RNN\_assignmentnet.ipynb - Colaboratory

```
# Tokenization
pos_review = tokenizer.texts_to_sequences(pos_review)

# padding
pos_review = pad_sequences(pos_review,maxlen=100,padding='post')

# prediction
review_predict = (model.predict(pos_review)>0.5).astype('int')

if review_predict[0] == 0:
    print("It's a negative review")
else:
    print("It's a positive review")

1/1 [=====] - 0s 53ms/step
It's a positive review
```